



# Original prompt (temp 0.7) vs TRS prompt (temp 0, greedy) — seven non-reasoning models


Mistral Nemo 12B (5-10B)




Qwen2.5 7B (5-10B)




Qwen3 8B (off) (5-10B)




Gemma 3 12B (10-20B)




Mistral Small 24B (20-50B)



Qwen2.5 32B (20-50B)



Gemma 3 27B (20-50B)



● original prompt · temp 0.7  
● TRS prompt · temp 0 (greedy)

Each panel is zoomed to its own accuracy range to reveal small setup differences. n = 1,000 paired items per condition per model.

Gemma 12B & 27B: original slightly higher. Qwen 7B: TRS clearly higher. Most cells nearly identical.